



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

---

Experiment No. 1
Analyze the Boston Housing dataset and apply appropriate Regression Technique
Date of Performance: 27-07-23
Date of Submission: 17-08-23



# Vidyavardhini's College of Engineering & Technology

## Department of Computer Engineering


---

**Aim:** Analyze the Boston Housing dataset and apply appropriate Regression Technique.

**Objective:** Ability to perform various feature engineering tasks, apply linear regression on the given dataset and minimise the error.

### Theory:

Linear Regression is a machine learning algorithm based on supervised learning. It performs a regression task. Regression models a target prediction value based on independent variables. It is mostly used for finding out the relationship between variables and forecasting. Different regression models differ based on – the kind of relationship between dependent and independent variables they are considering, and the number of independent variables getting used.



Linear regression performs the task to predict a dependent variable value (y) based on a given independent variable (x). So, this regression technique finds out a linear relationship between x (input) and y(output). Hence, the name is Linear Regression.

In the figure above, X (input) is the work experience and Y (output) is the salary of a person.

The regression line is the best fit line for our model.



# Vidyavardhini's College of Engineering & Technology

## Department of Computer Engineering

---

### **Dataset:**

The Boston Housing Dataset

The Boston Housing Dataset is derived from information collected by the U.S. Census Service concerning housing in the area of Boston MA. The following describes the dataset columns:

CRIM - per capita crime rate by town

ZN - proportion of residential land zoned for lots over 25,000 sq.ft.

INDUS - proportion of non-retail business acres per town.

CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)

NOX - nitric oxides concentration (parts per 10 million)

RM - average number of rooms per dwelling

AGE - proportion of owner-occupied units built prior to 1940

DIS - weighted distances to five Boston employment centres

RAD - index of accessibility to radial highways

TAX - full-value property-tax rate per \$10,000

PTRATIO - pupil-teacher ratio by town

B -  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town

LSTAT - % lower status of the population

MEDV - Median value of owner-occupied homes in \$1000's

### **Code:**

### **Conclusion:**



# Vidyavardhini's College of Engineering & Technology

## Department of Computer Engineering

---

The selected set of features for estimating house prices includes several attributes that capture various aspects of the towns, which can influence the median home value. These attributes include CRIM, ZN, INDUS, CHAS, NOX, RM, AGE, DIS, RAD, TAX, PTRATIO, B, LSTAT.


Crime rate, proportion of residential land, nitric oxide concentration, average number of rooms, accessibility to highways, property-tax rate, etc., these factors influence home prices in different neighborhoods. The target variable, MEDV (Median Home Value), is the central focus of our prediction model; these features can predict the median value of homes.

The Mean Squared Error of 0.03112933398095344 indicates the average squared difference between predicted and actual house prices in the model. Comparing this value to the range of actual house prices determines if the model's predictions are accurate enough or not.

```
import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
data=pd.read_csv('/content/BostonHousing (1).csv')
```

```
pd.read_csv('/content/BostonHousing (1).csv')
```



	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.4
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows x 14 columns

```
column_names = ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
```

```
print(data.head(5))
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	
	b	lstat	medv									
0	396.90	4.98	24.0									
1	396.90	9.14	21.6									
2	392.83	4.03	34.7									
3	394.63	2.94	33.4									
4	396.90	5.33	36.2									

```
print(np.shape(data))
```

(506, 14)

```
print(data.describe())
```

	crim	zn	indus	chas	nox	rm	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	
	age	dis	rad	tax	ptratio	b	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	
min	2.900000	1.129000	1.000000	187.000000	12.600000	0.320000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	
	lstat	medv					
count	506.000000	506.000000					

```

mean    12.653063    22.532806
std      7.141062     9.197104
min      1.730000     5.000000
25%      6.950000    17.025000
50%     11.360000    21.200000
75%     16.955000    25.000000
max     37.970000    50.000000

```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   crim        506 non-null    float64
 1   zn           506 non-null    float64
 2   indus        506 non-null    float64
 3   chas         506 non-null    int64
 4   nox          506 non-null    float64
 5   rm           506 non-null    float64
 6   age          506 non-null    float64
 7   dis          506 non-null    float64
 8   rad          506 non-null    int64
 9   tax          506 non-null    int64
10  ptratio      506 non-null    float64
11  b            506 non-null    float64
12  lstat        506 non-null    float64
13  medv         506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB

```

```
linr = LinearRegression()
```

```
data['medv'] = np.log1p(data['medv'])
```

```

X = data.drop(['medv','b'], axis=1)
Y = data['medv']

```

```
print(X)
```

```

      crim    zn  indus  chas   nox    rm   age    dis  rad  tax  \
0   0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900   1  296
1   0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671   2  242
2   0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671   2  242
3   0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622   3  222
4   0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622   3  222
..    ...    ...    ...    ...    ...    ...    ...    ...  ...  ...
501  0.06263   0.0  11.93    0  0.573  6.593  69.1  2.4786   1  273
502  0.04527   0.0  11.93    0  0.573  6.120  76.7  2.2875   1  273
503  0.06076   0.0  11.93    0  0.573  6.976  91.0  2.1675   1  273
504  0.10959   0.0  11.93    0  0.573  6.794  89.3  2.3889   1  273
505  0.04741   0.0  11.93    0  0.573  6.030  80.8  2.5050   1  273

      ptratio  lstat
0          15.3   4.98
1          17.8   9.14
2          17.8   4.03
3          18.7   2.94
4          18.7   5.33
..    ...    ...
501       21.0   9.67
502       21.0   9.08
503       21.0   5.64
504       21.0   6.48
505       21.0   7.88

```

```
[506 rows x 12 columns]
```

```
print(Y)
```

```

0    3.218876
1    3.117950
2    3.575151
3    3.538057
4    3.616309
...
501   3.152736
502   3.072693
503   3.214868
504   3.135494
505   2.557227
Name: medv, Length: 506, dtype: float64

```

```
x_train, x_test, y_train, y_test = train_test_split(X, Y, random_state=42, test_size=0.3)
```

```
print("x_train shape:",x_train.shape)
print("x_test shape:",x_test.shape)
print("y_train shape:",x_train.shape)
print("y_train shape:",x_test.shape)
```

```
x_train shape: (354, 12)
x_test shape: (152, 12)
y_train shape: (354, 12)
y_train shape: (152, 12)
```

```
linr.fit(x_train, y_train)
```

```
LinearRegression
LinearRegression()
```

```
y_pred = linr.predict(x_test)
```

```
print(mean_squared_error(y_test, y_pred))
```

```
0.03112933398095344
```

```
plt.scatter(y_test,y_pred,c='blue')
plt.xlabel("value")
plt.ylabel("Predicted value")
plt.title("True value vs predicted value : Linear Regression")
plt.show()
```

