



Experiment No. 4
Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:07-09-23
Date of Submission:14-09-23



Aim: Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, apply Random Forest Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

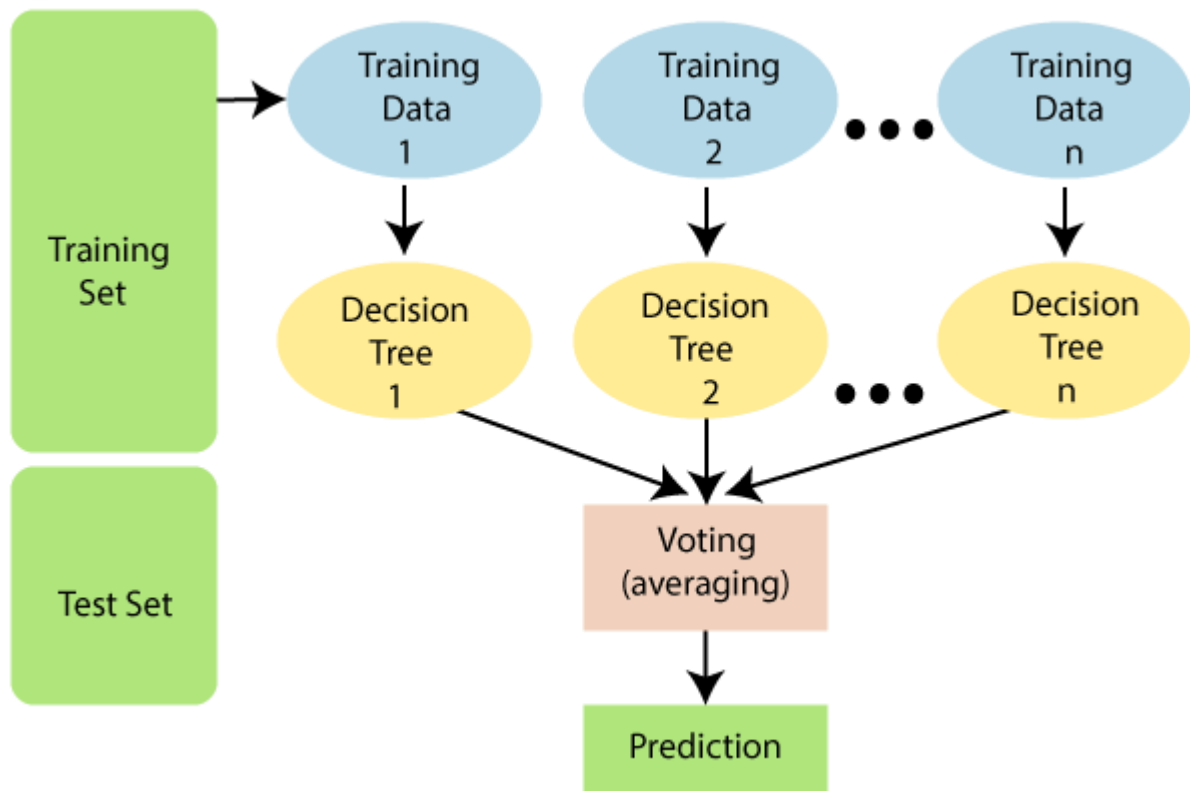
Theory:

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. It can be used for both Classification and Regression problems in ML. It is based on the concept of ensemble learning, which is a process of combining multiple classifiers to solve a complex problem and to improve the performance of the model.

As the name suggests, "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output.

The greater number of trees in the forest leads to higher accuracy and prevents the problem of overfitting.

The below diagram explains the working of the Random Forest algorithm:



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

The correlations among these variables are generally simpler in strength, they lack strong linear associations with one another.



Age exhibits a weak positive correlation with both education number and hours worked per week. Education numbers display a light positive correlation with capital gains. Also, there exists a weak negative correlation between capital gains and capital losses.

Accuracy: The model's accuracy is 85.44%, correctly predicting income levels for most instances.

Confusion Matrix: True positives (8015), False positives (628), and False negatives (1047), True negatives (1819) predictions.

Precision: Precision for income 0 = 0.08 and precision for income 1 = 0.74

Recall: Recall for income 0 = 0.93 and recall for income 1 = 0.63

F1-score: F1-score is the mean between precision and recall, indicating overall model effectiveness. It contains 0 for 0.91 and 1 for 0.68

Random Forest tends to provide better results than a Decision Tree. The Random Forest model combines the predictions of multiple Decision Trees, which can lead to improved accuracy and generalization.

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style='white', context='notebook', palette='deep')
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, learning_curve, train_test_split, KFold
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
# Adult dataset path
adult_dataset_path = "/content/adult.csv"
# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband

```
print ("Rows : " ,df.shape[0])
print ("Columns : " ,df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows : 48842
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'educational-num', 'marital-status', 'occupation', 'relationship', 'race', 'gender', 'income', 'capital-gain', 'capital-loss', 'hours-per-week', 'native-country']
```

```
Missing values : 0
```

```
Unique values :
age                74
workclass           9
fnlwgt            28523
education          16
educational-num    16
marital-status      7
occupation         15
relationship        6
race                5
gender              2
capital-gain       123
capital-loss        99
hours-per-week     96
native-country      42
income              2
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   age                  48842 non-null  int64
1   workclass            48842 non-null  object
2   fnlwgt              48842 non-null  int64
3   education            48842 non-null  object
4   educational-num      48842 non-null  int64
5   marital-status       48842 non-null  object
6   occupation           48842 non-null  object
7   relationship         48842 non-null  object
8   race                 48842 non-null  object
9   gender               48842 non-null  object
10  capital-gain         48842 non-null  int64
11  capital-loss         48842 non-null  int64
12  hours-per-week       48842 non-null  int64
13  native-country       48842 non-null  object
14  income               48842 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB

df.describe()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.422314
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.391414
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.000000

```
df.head()
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband
3	44	Private	160323	Some-col	10	Married-civ-spouse	Machine-op-inspct	Husband

```
# checking "?" total values present in particular 'workclass' feature
df_check_missing_workclass = (df['workclass']=='?').sum()
df_check_missing_workclass

2799

# checking "?" total values present in particular 'occupation' feature
df_check_missing_occupation = (df['occupation']=='?').sum()
df_check_missing_occupation

2809

# checking "?" values, how many are there in the whole dataset
df_missing = (df=='?').sum()
df_missing
```

```
age
workclass      2799
fnlwgt         0
education      0
educational-num 0
marital-status 0
occupation     2809
relationship   0
race           0
gender         0
capital-gain   0
capital-loss   0
hours-per-week 0
native-country 857
income         0
dtype: int64

percent_missing = (df=='?').sum() * 100/len(df)
percent_missing

age          0.000000
workclass    5.730724
fnlwgt       0.000000
education    0.000000
educational-num 0.000000
marital-status 0.000000
occupation   5.751198
relationship 0.000000
race         0.000000
gender       0.000000
capital-gain 0.000000
capital-loss 0.000000
hours-per-week 0.000000
native-country 1.754637
income       0.000000
dtype: float64

# find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()

age          48842
workclass    46043
fnlwgt       48842
education    48842
educational-num 48842
marital-status 48842
occupation   46033
relationship 48842
race         48842
gender       48842
capital-gain 48842
capital-loss 48842
hours-per-week 48842
native-country 47985
income       48842
dtype: int64

# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationshi
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-chil
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husban
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husban
						Married-		


```
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()
```

```
workclass      0
education      0
marital-status 0
occupation     10
relationship    0
race           0
gender         0
native-country 811
income         0
dtype: int64
```

```
from sklearn import preprocessing
# encode categorical variables using label Encoder
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()
```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country	i
0	Private	11th	Never-married	Machine-op-inspct	Own-child	Black	Male	United-States	
1	Private	HS-grad	Married-civ-spouse	Farming-fishing	Husband	White	Male	United-States	

```
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country	i
0	2	1	4	6	3	2	1	39	
1	2	11	2	4	0	4	1	39	
2	1	7	2	10	0	4	1	39	
3	2	15	2	6	0	2	1	39	

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marit sta
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	
2	28	336951	12	0	0	40	1	7	
3	44	160323	10	7688	0	40	2	15	

```
# look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 46033 non-null  int64
1   fnlwgt              46033 non-null  int64
2   educational-num     46033 non-null  int64
3   capital-gain        46033 non-null  int64
```

```

4 capital-loss 46033 non-null int64
5 hours-per-week 46033 non-null int64
6 workclass 46033 non-null int64
7 education 46033 non-null int64
8 marital-status 46033 non-null int64
9 occupation 46033 non-null int64
10 relationship 46033 non-null int64
11 race 46033 non-null int64
12 gender 46033 non-null int64
13 native-country 46033 non-null int64
14 income 46033 non-null int64

```

```


dtypes: int64(15)
memory usage: 5.6 MB

```

```

plt.figure(figsize=(14,10))
sns.heatmap(df.corr(),annot=True,fmt='.2f')
plt.show()

```



```

# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df info again whether everything is in right format or not
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    46033 non-null  int64
1   fnlwgt                 46033 non-null  int64
2   educational-num        46033 non-null  int64
3   capital-gain           46033 non-null  int64
4   capital-loss           46033 non-null  int64
5   hours-per-week         46033 non-null  int64
6   workclass              46033 non-null  int64
7   education              46033 non-null  int64
8   marital-status         46033 non-null  int64
9   occupation             46033 non-null  int64
10  relationship            46033 non-null  int64
11  race                   46033 non-null  int64
12  gender                 46033 non-null  int64
13  native-country         46033 non-null  int64
14  income                 46033 non-null  category

```

```
dtypes: category(1), int64(14)
memory usage: 5.3 MB
```

```
# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']
X.head(3)
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	maritsta
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	

```
y.head(3)
```

```
0    0
1    0
2    1
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y)
X_train.head()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	m
13554	58	196502	10	0	0	60	2	15	
46282	27	297457	9	0	0	40	2	11	
25679	27	30244	9	0	0	80	4	11	
8775	42	165309	9	0	0	50	2	11	

```
test_size = 0.20
seed = 7
num_folds = 10
scoring = 'accuracy'
# Params for Random Forest
num_trees = 100
max_features = 3
```

```
random_forest = RandomForestClassifier(n_estimators=250,max_features=5)
random_forest.fit(X_train, y_train)
predictions = random_forest.predict(X_test)
print("Accuracy: %s%" % (100*accuracy_score(y_test, predictions)))
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
```

```
Accuracy: 85.44617256060475%
[[8015  628]
 [1047 1819]]
      precision    recall  f1-score   support

     0       0.88      0.93      0.91      8643
     1       0.74      0.63      0.68      2866

 accuracy      0.85      0.85      0.85      11509
 macro avg      0.81      0.78      0.80      11509
 weighted avg      0.85      0.85      0.85      11509
```

