



Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:05/10/23
Date of Submission: 12/10/23



**Aim:** Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

**Theory:**

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

**Input:**

- $D$ , a set of  $d$  class labelled training tuples
- $k$ , the number of rounds (one classifier is generated per round)
- a classification learning scheme

**Output:** A composite model

**Method**

1. Initialize the weight of each tuple in  $D$  is  $1/d$
2. For  $i=1$  to  $k$  do // for each round
3. Sample  $D$  with replacement according to the tuple weights to obtain  $D_i$
4. Use training set  $D_i$  to derive a model  $M_i$
5. Compute  $\text{error}(M_i)$ , the error rate of  $M_i$
6.  $\text{Error}(M_i) = \sum w_j \cdot \text{err}(X_j)$
7. If  $\text{Error}(M_i) > 0.5$  then
8. Go back to step 3 and try again
9. endif
10. for each tuple in  $D_i$  that was correctly classified do



11. Multiply the weight of the tuple by  $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for

#### To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0
2. for  $i=1$  to  $k$  do // for each classifier
3.  $w_i = \log((1-\text{error}(M_i))/\text{error}(M_i))$  // weight of the classifiers vote
4.  $C = M_i(X)$  // get class prediction for X from  $M_i$
5. Add  $w_i$  to weight for class C
6. end for
7. Return the class with the largest weight.

#### Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.



race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

#### **Code:**

#### **Conclusion:**

Accuracy: 0.865, indicating that the model correctly predicts the income level.

Confusion Matrix: True Positives (637), True Negatives (144), False Positives (379), False Negatives(854) predictions.

Precision: The precision for income 0 = 0.88 and the precision for income 1 = 0.79.

Recall: The recall for income 0 = 0.94 and recall for income 1 = 0.63

F1-Score: The F1-score is the mean between precision and recall, indicating overall model effectiveness. It contains 0 for 0.91 and 1 for 0.70

Both Random Forest and AdaBoost are capable of delivering high accuracy and exhibit reduced susceptibility to overfitting. Nevertheless, Random Forest tends to be more resilient to variations in hyperparameter tuning, often requiring less extensive adjustments. Additionally, Random Forest offers the advantage of feature importance analysis, enhancing its interpretability, whereas AdaBoost's sequential nature can result in a lower level of interpretability. In cases involving imbalanced data, AdaBoost outperforms Random Forest by assigning greater weights to minority class samples, thus addressing the class imbalance issue more effectively. To summarize, AdaBoost and Random Forest represent potent ensemble algorithms, but their performance is subject to variation based on hyperparameter settings and the characteristics of the dataset.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import io
from sklearn.metrics import accuracy_score, precision_score, f1_score, confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
```

```
import os
for dirname, _, filenames in os.walk('/content/adult.csv'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
file = ('/content/adult.csv')
df = pd.read_csv(file)
```

```
print(df.head())
```

```

age workclass  fnlwgt      education  education.num marital.status \
0    90      ?    77053      HS-grad           9      Widowed
1    82  Private  132870      HS-grad           9      Widowed
2    66      ?   186061  Some-college          10      Widowed
3    54  Private  140359      7th-8th           4      Divorced
4    41  Private  264663  Some-college          10      Separated

      occupation  relationship    race    sex  capital.gain \
0              ?  Not-in-family  White  Female           0
1  Exec-managerial  Not-in-family  White  Female           0
2              ?      Unmarried  Black  Female           0
3  Machine-op-inspct  Unmarried  White  Female           0
4   Prof-specialty    Own-child  White  Female           0

capital.loss  hours.per.week  native.country  income
0         4356              40  United-States  <=50K
1         4356              18  United-States  <=50K
2         4356              40  United-States  <=50K
3         3900              40  United-States  <=50K
4         3900              40  United-States  <=50K
```

```
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
None
```

```

for i in df.columns:
    t = df[i].value_counts()
    index = list(t.index)
    print("Count of ? in", i)
    for i in index:
        temp = 0
        if i == '?':
            print(t['?'])
            temp = 1
            break
    if temp == 0:
        print("0")
```

Streaming output truncated to the last 5000 lines.

0

	age	workclass	fnlwtg	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th		Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	

```
6 38 Private 150601 10th 6 Separated
```

```

      occupation relationship race sex capital.gain \
1  Exec-managerial Not-in-family White Female 0
3  Machine-op-inspct Unmarried White Female 0
4  Prof-specialty Own-child White Female 0
5  Other-service Unmarried White Female 0
6  Adm-clerical Unmarried White Male 0

```

```

capital.loss hours.per.week native.country income
1      4356      18 United-States 0
3      3900      40 United-States 0
4      3900      40 United-States 0
5      3770      45 United-States 0
6      3770      40 United-States 0

```

```
df_more=df.loc[df['income'] == 1]
print(df_more.head())
```

```

      age workclass fnlwgt education education.num marital.status \
7      74 State-gov 88638 Doctorate 16 Never-married
10     45 Private 172274 Doctorate 16 Divorced
11     38 Self-emp-not-inc 164526 Prof-school 15 Never-married
12     52 Private 129177 Bachelors 13 Widowed
13     32 Private 136204 Masters 14 Separated

```

```

      occupation relationship race sex capital.gain \
7  Prof-specialty Other-relative White Female 0
10 Prof-specialty Unmarried Black Female 0
11 Prof-specialty Not-in-family White Male 0
12 Other-service Not-in-family White Female 0
13 Exec-managerial Not-in-family White Male 0

```

```

capital.loss hours.per.week native.country income
7      3683      20 United-States 1
10     3004      35 United-States 1
11     2824      45 United-States 1
12     2824      20 United-States 1
13     2824      55 United-States 1

```


```
workclass_types = df_more['workclass'].value_counts()
labels = list(workclass_types.index)
aggregate = list(workclass_types)
print(workclass_types)
print(aggregate)
print(labels)
```

```


Private      4876
Self-emp-not-inc 714
Local-gov    609
Self-emp-inc 600
Federal-gov  365
State-gov    344
Name: workclass, dtype: int64
[4876, 714, 609, 600, 365, 344]
['Private', 'Self-emp-not-inc', 'Local-gov', 'Self-emp-inc', 'Federal-gov', 'State-gov']

```

```
plt.figure(figsize=(7,7))
plt.pie(aggregate, labels=labels, autopct='%1.1f%%', shadow = True)
plt.axis('equal')
plt.show()
```



```
#Count plot on single categorical variable
sns.countplot(x='income', data = df)
plt.show()
df['income'].value_counts()
```




```
0    22661
1     7508
Name: income, dtype: int64
```

```
#Plot figsize
plt.figure(figsize=(10,7))
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
print(plt.show())
plt.figure(figsize=(10,7))
sns.distplot(df['age'], color="red", bins=100)
plt.ylabel("Distribution", fontsize = 10)
plt.xlabel("Age", fontsize = 10)
plt.show()
```



```
<ipython-input-14-c01c35a847eb>:3: FutureWarning: The default value of numeric_only
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```



```
#To find distribution of categorical columns w.r.t income
```

```
fig, axes = plt.subplots(figsize=(20, 10))
```

```
plt.subplot(231)
```

```
sns.countplot(x='workclass',
```

```
hue='income',
```

```
data = df,
```

```
palette="BuPu")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(232)
```

```
sns.countplot(x='marital.status',
```

```
hue='income',
```

```
data = df,
```

```
palette="deep")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(233)
```

```
sns.countplot(x='education',
```

```
hue='income',
```

```
data = df,
```

```
palette = "autumn")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(234)
```

```
sns.countplot(x='relationship',
```

```
hue='income',
```

```
data = df,
```

```
palette = "inferno")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(235)
```

```
sns.countplot(x='sex',
```

```
hue='income',
```

```
data = df,
```

```
palette = "coolwarm")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(236)
```

```
sns.countplot(x='race',
```

```
hue='income',
```

```
data = df,
```


```
palette = "cool")
```

```
plt.xticks(rotation=90)
```

```
plt.subplots_adjust(hspace=1)
```

```
plt.show()
```

```
<ipython-input-16-e1b6d1f0108f>:3: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed in a future version.
plt.subplot(231)
```



```
df1 = df.copy()
categorical_features = list(df1.select_dtypes(include=['object']).columns)
print(categorical_features)
df1
```

```
['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gai
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Female	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	
5	34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	
6	38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	
...	...	...	...	...	...	...	...	...	...	...	...
32556	22	Private	310152	Some-college	10	Never-married	Protective-serv	Not-in-family	White	Male	
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	

30169 rows × 15 columns

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for feat in categorical_features:
    df1[feat] = le.fit_transform(df1[feat].astype(str))
df1

X = df1.drop(columns = ['income'])
y = df1['income'].values
# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
print ("Train set size: ", X_train.shape)
print ("Test set size: ", X_test.shape)

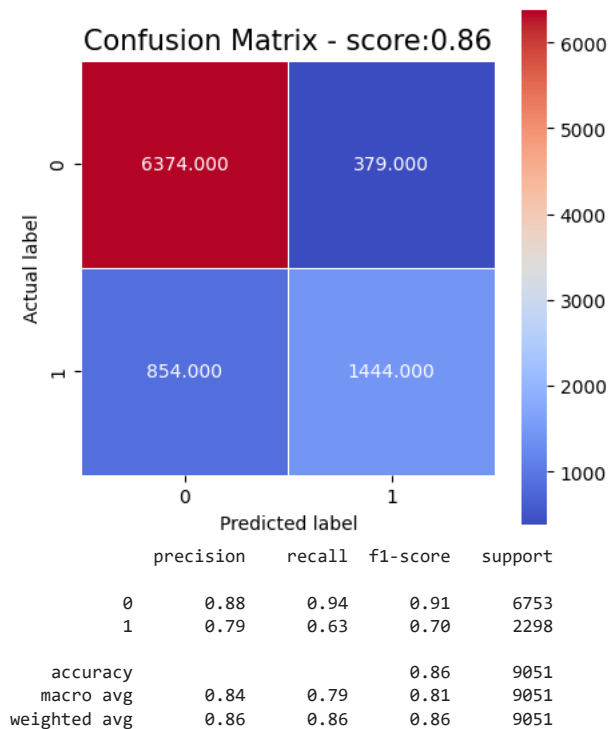
from sklearn.ensemble import AdaBoostClassifier
# Train Adaboost Classifier
abc = AdaBoostClassifier(n_estimators = 300, learning_rate=1)
abc_model = abc.fit(X_train, y_train)
#Prediction
y_pred_abc = abc_model.predict(X_test)

Train set size: (21118, 14)
Test set size: (9051, 14)

print("Accuracy: ", accuracy_score(y_test, y_pred_abc))
print("F1 score :", f1_score(y_test, y_pred_abc, average='binary'))
print("Precision : ", precision_score(y_test, y_pred_abc))
```

Accuracy: 0.8637719588995691  
 F1 score : 0.7008007765105557  
 Precision : 0.7921009325287987

```
cm = confusion_matrix(y_test, y_pred_abc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap = "coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,y_pred_abc), 2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_abc))
```




```
from sklearn.ensemble import GradientBoostingClassifier
#Training the model with gradient boosting
gbc = GradientBoostingClassifier(
    learning_rate = 0.1,
    n_estimators = 500,
    max_depth = 5,
    subsample = 0.9,
    min_samples_split = 100,
    max_features='sqrt',
    random_state=10)
gbc.fit(X_train,y_train)
# Predictions
y_pred_gbc = gbc.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_gbc))
print("F1 score : ", f1_score(y_test, y_pred_gbc, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_gbc))
```

Accuracy : 0.8689647552756602  
 F1 score : 0.7218574108818011  
 Precision : 0.7828077314343845

```
rms = np.sqrt(mean_squared_error(y_test, y_pred_gbc))
print("RMSE for gradient boost: ", rms)
```

```
cm = confusion_matrix(y_test, y_pred_gbc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot = True, fmt=".3f", linewidths = 0.5, square = True, cmap= "coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,y_pred_gbc),2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_gbc))
```

RMSE for gradient boost: 0.3619879068758235



	precision	recall	f1-score	support
0	0.89	0.94	0.91	6753
1	0.78	0.67	0.72	2298
accuracy	0.84	0.80	0.82	9051


```
import xgboost as xgb
from xgboost import XGBClassifier
#Training the model with gradient boosting
xgboost = XGBClassifier(learning_rate=0.01,
    colsample_bytree = 0.4,
    n_estimators=1000,
    max_depth=20,
    gamma=1)
xgboost_model = xgboost.fit(X_train, y_train)

# Predictions
y_pred_xgboost = xgboost_model.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_xgboost))
print("F1 score : ", f1_score(y_test, y_pred_xgboost, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_xgboost))

rms = np.sqrt(mean_squared_error(y_test, y_pred_xgboost))
print("RMSE for xgboost: ", rms)

cm = confusion_matrix(y_test, y_pred_xgboost)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap ="coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:'+str(round(accuracy_score(y_test,y_pred_xgboost),2)), size = 15);
plt.show()
print(classification_report(y_test,y_pred_xgboost))
```

Accuracy : 0.8655397193680257  
F1 score : 0.7090604829070045  
Precision : 0.786737400530504  
RMSE for xgboost: 0.3666882608319693



```
from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, xgboost.predict_proba(X_test)[:,:1])
plt.figure(figsize = (10,5))
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC CURVE for Xgboost')
plt.show()
```

