



Experiment No. 7
Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model
Date of Performance: 12/10/23
Date of Submission: 16/10/23



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

**Aim:** Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model.

**Objective:** Able to perform various feature engineering tasks, perform dimensionality reduction on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

#### **Theory:**

In machine learning classification problems, there are often too many factors on the basis of which the final classification is done. These factors are basically variables called features. The higher the number of features, the harder it gets to visualize the training set and then work on it. Sometimes, most of these features are correlated, and hence redundant. This is where dimensionality reduction algorithms come into play. Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables. It can be divided into feature selection and feature extraction.

#### **Dataset:**

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.



## Vidyavardhini's College of Engineering & Technology

### Department of Computer Engineering

---

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

#### **Code:**

#### **Conclusion:**

Accuracy: After undergoing dimensionality reduction, it demonstrates an accuracy of around 0.821.

Precision: For the  $\leq 50K$  class, the model exhibits a precision of 0.84 and for the  $> 50K$  class, the model exhibits a precision of 0.72

Recall: For the  $\leq 50K$  class, the model exhibits a recall of 0.95, and for the  $> 50K$  class, the model exhibits a recall of 0.43

F1-score: For the  $\leq 50K$  class, the model exhibits a F1-score of 0.89 and for the  $> 50K$  class, the model exhibits a F1-score of 0.54.

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
```

```
for dirname, _, filenames in os.walk('/content/adult (1).csv'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
df=pd.read_csv("/content/adult (1).csv")
```

```
df.head
```

```
<bound method NDFrame.head of
0      90      ?  77053      HS-grad      9      Widowed
1      82 Private 132870      HS-grad      9      Widowed
2      66      ? 186061 Some-college     10      Widowed
3      54 Private 140359      7th-8th      4      Divorced
4      41 Private 264663 Some-college     10      Separated
...      ...      ...      ...      ...      ...
32556  22 Private 310152 Some-college     10      Never-married
32557  27 Private 257302 Assoc-acdm     12 Married-civ-spouse
32558  40 Private 154374      HS-grad      9 Married-civ-spouse
32559  58 Private 151910      HS-grad      9      Widowed
32560  22 Private 201490      HS-grad      9      Never-married
```

```
      occupation relationship race sex capital.gain \
0      ? Not-in-family White Female      0
1  Exec-managerial Not-in-family White Female      0
2      ? Unmarried Black Female      0
3  Machine-op-inspct Unmarried White Female      0
4  Prof-specialty Own-child White Female      0
...      ...      ...      ...      ...
32556  Protective-serv Not-in-family White Male      0
32557  Tech-support Wife White Female      0
32558  Machine-op-inspct Husband White Male      0
32559  Adm-clerical Unmarried White Female      0
32560  Adm-clerical Own-child White Male      0
```

```
      capital.loss hours.per.week native.country income
0      4356      40 United-States <=50K
1      4356      18 United-States <=50K
2      4356      40 United-States <=50K
3      3900      40 United-States <=50K
4      3900      40 United-States <=50K
...      ...      ...      ...
32556      0      40 United-States <=50K
32557      0      38 United-States <=50K
32558      0      40 United-States >50K
32559      0      40 United-States <=50K
32560      0      20 United-States <=50K
```

```
[32561 rows x 15 columns]>
```

```
df.columns
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education.num',
      'marital.status', 'occupation', 'relationship', 'race', 'sex',
      'capital.gain', 'capital.loss', 'hours.per.week', 'native.country',
      'income'],
      dtype='object')
```

```
df.shape
```

```
(32561, 15)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
```

```

10 capital.gain    32561 non-null int64
11 capital.loss    32561 non-null int64
12 hours.per.week  32561 non-null int64
13 native.country  32561 non-null object
14 income          32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
df[df == '?'] = np.nan
```

```
df.isnull().sum()
```

```

age                0
workclass          1836
fnlwgt             0
education           0
education.num       0
marital.status      0
occupation          1843
relationship        0
race               0
sex               0
capital.gain        0
capital.loss        0
hours.per.week      0
native.country      583
income             0
dtype: int64

```

```

for col in ['workclass', 'occupation', 'native.country']:
    df[col].fillna(df[col].mode()[0], inplace=True)
df.isnull().sum()

```

```

age                0
workclass          0
fnlwgt             0
education           0
education.num       0
marital.status      0
occupation          0
relationship        0
race               0
sex               0
capital.gain        0
capital.loss        0
hours.per.week      0
native.country      0
income             0
dtype: int64

```

```

# converting categorical Columns
df.replace({'Sex':{'male':0,'female':1}, 'Embarked':{'S':0,'C':1,'Q':2}}, inplace=True)
X = df.drop(['income'], axis=1)
y = df['income']
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
from sklearn import preprocessing
categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])

```

```

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
X_train.head()

```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gai
0	0.101484	2.600478	-1.494279	-0.332263	1.133894	-0.402341	-0.782234	2.214196	0.39298	-1.430470	-0.14518
1	0.028248	-1.884720	0.438778	0.184396	-0.423425	-0.402341	-0.026696	-0.899410	0.39298	0.699071	-0.14518
2	0.247956	-0.090641	0.045292	1.217715	-0.034095	0.926666	-0.782234	-0.276689	0.39298	-1.430470	-0.14518
3	-0.850587	-1.884720	0.793152	0.184396	-0.423425	0.926666	-0.530388	0.968753	0.39298	0.699071	-0.14518
4	-0.044989	-2.781760	-0.853275	0.442726	1.523223	-0.402341	-0.782234	-0.899410	0.39298	0.699071	-0.14518

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

LR = LogisticRegression()
LR.fit(X_train, y_train)
y_pred = LR.predict(X_test)
accuracy_score(y_test, y_pred)

0.8216808271061521

from sklearn.decomposition import PCA
pca = PCA()
X_train = pca.fit_transform(X_train)
pca.explained_variance_ratio_

X = df.drop(['income'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
pca= PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
dim = np.argmax(cumsum >= 0.90) + 1
print('The number of dimensions required to preserve 90% of variance is',dim)

The number of dimensions required to preserve 90% of variance is 12

X = df.drop(['income', 'native.country', 'hours.per.week'], axis=1)
y = df['income']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex']
for feature in categorical:
    label = preprocessing.LabelEncoder()
    X_train[feature] = label.fit_transform(X_train[feature])
    X_test[feature] = label.transform(X_test[feature])
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

LR2 = LogisticRegression()
LR2.fit(X_train, y_train)



▾ LogisticRegression  

    LogisticRegression()



y_pred = LR2.predict(X_test)
accuracy_score(y_test, y_pred)

0.8227044733340158

from sklearn.metrics import confusion_matrix
import pandas as pd
confusion = confusion_matrix(y_test, y_pred)
df_confusion = pd.DataFrame(confusion, columns=['Predicted No', 'Predicted Yes'], index=['Actual No', 'Actual Yes'])
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
<=50K	0.84	0.95	0.89	7410
>50K	0.72	0.43	0.54	2359
accuracy			0.82	9769
macro avg	0.78	0.69	0.72	9769
weighted avg	0.81	0.82	0.81	9769

